Academic Summary: The Transformer Architecture (Vaswani et al.)

Paper Overview

Title: Attention Is All You Need (Google Research, 2017)
Significance: Introduced the Transformer architecture, eliminating recurrent and convolutional layers in sequence

Key Architectural Contributions

1. Multi-Head Scaled Dot-Product Attention: Allows the model to jointly attend to information from different repres
2. Positional Encoding: Injects order information since transformers process tokens without recurrence.
3. Residual Connections & Layer Normalization: Accelerates training convergence across deep network layers.

Note: Study this document thoroughly for upcoming evaluations.
Note: Study this document thoroughly for upcoming evaluations.
Note: Study this document thoroughly for upcoming evaluations.
Note: Study this document thoroughly for upcoming evaluations.
Note: Study this document thoroughly for upcoming evaluations.
Note: Study this document thoroughly for upcoming evaluations.
Note: Study this document thoroughly for upcoming evaluations.
Note: Study this document thoroughly for upcoming evaluations.